

CSE312 - Operating Systems - HW#4

Fatih Kaan Salgır - 171044009

Design Explanation




Figure 1: Disk block alignments according to different block sizes.

Since there are lots of data blocks, data block area in the figure represents only 1% of the whole data blocks in the disk.

Creating an empty file system

To test `makeFileSystem`, disk file is examined with `hexdump` command in linux. First part is the contents of the super block. The second is Free Block Bitmap is the first blocks of the bitmap is marked as full, because root directory is marked as full in bitmap. And linked one after another in FAT Table.

```
> ./makeFileSystem 4 mySystem.dat 86 du -b mySystem.dat 86 hexdump mySystem.dat
16777216 mySystem.dat
00000000 1000 0000 1000 0000 1000 0000 1000 0000
00000010 0fe0 0000 2000 0000 0002 0000 2000 0001
00000020 e000 0000 0000 0002 0000 00fe 0000 0000
00000030 0000 0000 0000 0000 0000 0000 0000 0000
*
00010000 ffff 0000 0000 0000 0000 0000 0000 0000
00010100 0000 0000 0000 0000 0000 0000 0000 0000
*
00020000 0fff 0fff 0003 0004 0005 0006 0007 0008
00020100 0009 000a 000b 000c 000d 000e 000f 0fff
00020200 0000 0000 0000 0000 0000 0000 0000 0000
*
10000000
```

Exactly 16MB for 4KB block size

Super Block

Free Block Bitmap

FAT Table

16777216 in hexadecimal format

Figure 2: Testing `makeFileSystem` with 4KB block size

Super Block

Super block holds information about;

- block size,
- number of blocks,
- where the free bitmap, fat table root directory and data blocks starts and their sizes
- number of free blocks

```
// filesystem.h
struct super_block {
    uint32_t block_size;
    uint32_t block_count;
    uint32_t free_bitmap_start;
    uint32_t free_block_size;
    uint32_t no_free_blocks;
    uint32_t fat_table_start;
    uint32_t fat_table_block_size;
    uint32_t root_start;
    uint32_t root_size;
    uint32_t data_start;
    uint32_t data_size;
};
```

Free Blocks

There are about 2^{12} data blocks (some of the blocks are reserved). In bitmap representation every data block is mapped to a bit. Therefore there must be 2^{12} bits in bitmap.

$$2^{12} \text{ bits} = \frac{2^{12}}{2^3} \text{ bytes} = 512 \text{ bytes}$$

The bitmap fits inside a block, because the minimum block size is 512 bytes. During implementation of the filesystem bit manipulation is needed. Bitmap implemented as `uint32_t` array with size of 128 which is 2^{12} bits.

FAT Table

There are 2^{12} entries and every entry is 12 bits. Therefore the table size;

$$\frac{2^{12} \times 12}{2^3} = 6 \times 2^{10} \text{ Bytes} = 6 \text{ KB}$$

So the number of blocks needed for FAT table depends on the block size. If the block size is 512 Bytes, then 12 blocks needed. On the other hand if the block size is 4 KB, then 1.5 blocks will be enough. To align properly numbers completed to integers. Final table is;

Table 1: Number of blocks needed for FAT table, according to different block size

Block Size	FAT-12	No blocks needed
0.5 KB	2MB	12
1 KB	4MB	6
2 KB	8MB	3
4 KB	16MB	2

```
// filesystem.h
struct uint12 {
    uint16_t x : 12;
};
typedef struct uint12 uint12;
extern uint12 fat_table[NO_BLOCKS];
```

Directory Entry

32 bytes used for each entry. Address is the block number. Attributes shows if the entry is file or directory. If the entry is a file then size shows how many bytes the file size. However if the entry is a directory, the size shows how many files/folders inside the directory.

Size of the directory can be represented with 3 bytes, since it doesn't exceed maximum file size in the file system. In order to store the size in 3 bytes, array with size 3 and `uint8_t` type is used. `pack3b()` and `unpack3b()` functions are used to fit 32 bit integer into 24 bits.




Figure 3: Directory entries - number of bytes used

```
struct dir_entry {
    char name[20];           // 20 bytes
    uint16_t address;        // 2 bytes FAT-12 (used only least significant 12 bits)
    uint8_t attr;            // 1 byte
    uint8_t size[3];         // 3 bytes : max 16 MB size
    uint8_t time[3];         // 3 bytes : HMS
    uint8_t date[3];         // 3 bytes : DMY (Y: 1900+)
} __attribute__((packed)); // required for struct to be exactly 32 bytes
```

Running & Test Results

File system manipulation functions are implemented in `fileSystemOper.c`.

Table 2: Function names for corresponding operations

Operation	Function
dir	<code>fs_dir()</code>
mkdir	<code>fs_mkdir()</code>
rmdir	<code>fs_rmdir()</code>
dumpe2fs	<code>dumpe2fs()</code>
write	<code>fs_write()</code>
read	<code>fs_read()</code>
del	<code>fs_del()</code>

Test cases in the assignment pdf

```
> ./makeFileSystem 4 mySystem.data
> ./fileSystemOper mySystem.data mkdir '\usr'
> ./fileSystemOper mySystem.data mkdir '\usr\ysa'
> ./fileSystemOper mySystem.data mkdir '\bin\ysa'
ERROR: Path is incorrect, couldn't find: bin
> echo contents of linuxFile.data > linuxFile.data
> ./fileSystemOper mySystem.data write '\usr\file2' linuxFile.data
> ./fileSystemOper mySystem.data write '\file3' linuxFile.data
> ./fileSystemOper mySystem.data dir '\'
25-05-2021 18:07 <DIR>      usr
25-05-2021 18:07      27 file3
      1 File(s)      27 bytes
      1 Dir(s) 16625664 bytes free
> ./fileSystemOper mySystem.data del '\usr\ysa\file1'
> ./fileSystemOper mySystem.data dump2fs
==== SUPERBLOCK ====
Block Size:      4096
Block Count:     4096
Free Bitmap Start: 4096
Free Blocks Size: 4096
Number of Free Blocks: 4060
FAT Table Start: 8192
FAT Table Block Size: 2
Root Start:      73728
Root Size:       57344
Data Start:      131072
Data Size:       16646144
===== BITMAP =====
  0 : 1111111111111111 1101100000000000 0000000000000000 0000000000000000
 64 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
128 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
192 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
256 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
320 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
384 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
448 : 0000000000000000 0000000000000000 0000000000000000 0000000000000000
...(truncated)
===== OCCUPIED BLOCK - FILENAME =====
16 - usr
17 - ysa
19 - file2
20 - file3
Total number of directories: 2
Total number of files: 2
> ./fileSystemOper mySystem.data read '\usr\file2' linuxFile2.data
> cmp linuxFile2.data linuxFile.data
```